

Karthikeyan Devadoss

Berlin, Germany · itskarthiksid@yahoo.com · +49 1521 0924553

<https://job-boards.greenhouse.io/databricks/jobs/8648494002>

Dear Hiring Team,

Databricks sits at the centre of how serious teams put data and AI into production. Forward Deployed Engineer in Berlin means owning architecture and delivery with customers — the customer-facing builder path I am deliberately moving toward. I am writing about the Forward Deployed Engineer role.

What I bring

- About ten years of Java/Spring backend work, including US large-enterprise delivery: Marsh in Phoenix (Bluestream / Digital Broker for Affinity — ~95,000 colleagues across 130 countries) and BCBSA in Chicago (national hub connecting state Blue Cross Blue Shield plans, ~118 million members system-wide).
- Since early 2026, one of a handful of engineers in an initial core AI team at a large retail-energy company: production AI triage with Azure OpenAI, vector retrieval, and a human approval step before any action, now used by several teams.
- Own agentic software-delivery platform (live): approval-gated agents, RAG with labelled retrieval evals, MCP tool boundaries — public proof that AI systems work can hold up past a demo.

How it fits this role

- Took part in architecture discussions and design decisions with the lead; comfortable owning design choices on production systems.
- Backend depth (Java/Spring, APIs, cloud) plus hands-on AI triage with retrieval and approval gates.
- Berlin-listed role; English; experience delivering for large US enterprises (Marsh Phoenix, BCBSA Chicago).

Lakehouse and AI platforms only matter when customer productions stay up. I want the seat where architecture decisions meet measurable delivery outcomes. At this point in my career I want that customer-facing, architecture-adjacent seat more than another pure backend lane.

If Databricks is how customers turn data into production AI, I want to be the FDE helping them get there in Berlin.

Kind regards,

Karthikeyan Devadoss